

When Agentic Digital Twin Telemetry Measures Itself: A Provenance Audit and a Gate

Umberto Altieri

Independent Researcher

altieriumb7@gmail.com

Abstract—Agentic Digital Twins delegate diagnostic and maintenance decisions to tool-augmented language models, so predicting when such an agent is about to fail—and abstaining—is a safety requirement. The signals proposed for this are token-level uncertainty, verbalized confidence, and tool-execution telemetry. We audited a working pipeline built to evaluate exactly those signals on an industrial prognostics benchmark, and found that none of its features measured the system it claimed to describe. Token logprobs were synthesised by the logging layer from a hard-coded constant; verbalized confidence was never elicited; the tool-error rate was inferred by substring-matching observation text; the grading harness wrote its verdict into that same text; the telemetry itself was re-derived by parsing run logs, inflating every structural feature; and on 30 of 75 benchmark scenarios the tool layer fabricates the analytical outcome, so the label records orchestration rather than prognostic accuracy. Each defect is an ordinary engineering decision that produces a numerically healthy column. We report six such defects with the cheap test that detects each, release a gate that automates those tests and fails closed, and show what the same analysis yields once the pipeline is repaired: on a re-collected corpus of 225 trajectories with provider-returned logprobs, early-step entropy separates success from failure on two of three backbones (AUROC 0.776 and 0.746) while the model’s own post-hoc confidence does not (0.494), and the effect disappears entirely when backbones are pooled.

Index Terms—Digital Twin engineering, agentic systems, uncertainty quantification, data provenance, reproducibility, predictive maintenance.

I. INTRODUCTION

Industrial Digital Twins increasingly delegate work to LLM agents that orchestrate diagnostic tools over standard interfaces [1]. Benchmarks in this setting report that even strong backbones complete only around two thirds of tasks [2], [3]. Because acting on a wrong prognosis costs far more than deferring, a twin that recognises its own unreliable runs and abstains is worth more than one that is merely more accurate. A model is calibrated when its stated confidence matches empirical correctness [8]; for agents the analogous question is *trajectory calibration*, predicting run success from signals observable during execution of a ReAct-style loop [4]—token-level uncertainty [5], verbalized confidence [6], and execution telemetry [7].

This paper is not about whether that works. It is about a prior question we did not think to ask of our own pipeline: *are these features measurements?* For the pipeline we audited,

none of them were. The defects are not exotic—a fallback that keeps a long sweep running, a default for an unparsed field, a status message appended to a trace—and each yields a column that passes every check short of asking where the numbers came from.

We contribute (i) six documented defects with a cheap detector for each, (ii) a released gate that automates them and refuses a corpus that fails, and (iii) a demonstration of what changes when the same analysis runs on repaired instrumentation.

II. SETTING

We state the setup explicitly, because the defects below are only visible once one is precise about what is recorded and when.

A scenario is a natural-language maintenance request over a real asset dataset: predict remaining useful life for every unit in a turbofan fleet, classify a bearing fault, judge whether a maintenance window meets a safety policy. The agent answers by interleaving a *thought*, an *action* naming a tool, an *action input*, and the *observation* the tool returns, repeating until it emits a terminating action or exhausts a step budget. A grader then compares the final answer against scenario ground truth, yielding one binary label per run. That sequence of steps, with the label, is a *trajectory*.

Four families of features are extracted from it. *Token* features summarise the per-token log-probabilities of the agent’s thoughts—their mean, their variance, and the entropy of the first two steps. *Verbalized* confidence is the agent’s own stated certainty. *Structural* features describe the shape of the trace: how many steps, how many tool calls, how often an action repeats, whether it terminated properly. *Error-rate* features count tool calls that failed. The premise of trajectory calibration is that some combination of these predicts the label well enough to support abstention.

Everything in this paper follows from asking where each of those numbers came from.

A. Why the twin needs this

The asymmetry that makes calibration worth having is economic. Acting on a prognosis that is wrong—scheduling no intervention on a component about to fail—costs an unplanned outage. Deferring to a human engineer costs a review. The two

TABLE I

DEFECTS FOUND, THE EVIDENCE, AND THE TEST THAT DETECTS EACH. EVERY VALUE IN EVERY COLUMN WAS INDIVIDUALLY WELL-FORMED.

Defect	Evidence	Detector
Logprobs synthesised	All 203,107 tokens fit the logger's own fallback generator $\mathcal{N}(-0.40, 0.225^2)$; observed -0.4030 ± 0.2182	Compare moments against what the fallback would emit
Confidence never elicited	162/167 trajectories carry a hard-coded 0.75; parsed once per run, so its gradient is 0 by construction	Count distinct values; a constant cannot inform
Error rate is prose, not protocol	<code>status_code</code> set by matching "Error"/"failed" in observation text; no response metadata is read	Require the status the tool server returned
Grader verdict leaks	Terminal observation reads "Task finished successfully"/"Task failed"; the <i>only</i> observation for 71 single-step runs	Non-monotone learner far above a feature's univariate AUROC
Telemetry reconstructed	Re-derived by parsing logs; the single-use <code>Finish</code> appears 866 times, 790 mid-trajectory	Check trace shape against what the agent can emit
Label fabricated	<code>predict_rul</code> returns ground truth plus $\mathcal{N}(-3, 14^2)$ at fixed seed, giving MAE 9.01 inside the grader's [7, 18] window; <code>classify_faults</code> invents its ground truth too	Read the tool, not its output

differ by orders of magnitude, so a twin that abstains on its least reliable decisions converts the expensive error into the cheap one, even at mediocre discrimination.

That argument only holds if the confidence estimate is real. A twin abstaining on a signal that is actually a logging artifact does not reduce risk; it redistributes it arbitrarily while appearing instrumented, which is worse than having no abstention at all, because the operator now trusts a control that does nothing. This is why provenance is a safety property here rather than a matter of experimental hygiene.

III. SIX PROVENANCE DEFECTS

The audited pipeline evaluates five calibration schemes over 167 trajectories from seven backbones on PHMForge [2]. Illegitimate features are a leading cause of irreproducibility in ML-based science [10]; what follows is that failure mode arising from agent instrumentation rather than from dataset construction. Table I summarises what each of its feature families actually contained.

Fig. 1 shows the two clearest cases. The test that produced it is one line—compare the empirical moments of a generation-derived feature against what the fallback path would emit—and it is the cheapest guard we know of against an uncertainty feature that is not a measurement.




Fig. 1. Provenance evidence. **Left:** all 203,107 logged token logprobs against the density of the logging layer's own fallback generator at its default confidence. The fit is exact; the spike at zero is the generator's clipping. **Right:** verbalized confidence takes one value on 162 of 167 trajectories. Neither family was obtained from a model.

Three points generalise beyond this codebase. First, the logprob chain was broken in *four* independent places—the client never requested them, one of two agent classes discarded them, the default serialisation branch omitted them, and the logger synthesised a replacement. Fixing any three would have changed nothing observable, because the fourth still supplied plausible numbers. That is how a fabricated feature survives review: the defect is distributed across individually reasonable layers, and the fallback removes the error signal that would force someone to look.

Second, the last defect concerns the *target*, and no feature audit could have found it. Every column can be a perfect measurement while the quantity being predicted is not the one of interest. It is worth seeing directly. The tool that is supposed to run the trained model instead does this:

```
gt = load_rul_ground_truth(fd)
rng = np.random.default_rng(seed=42)
shifts = rng.normal(loc=-3.0, scale=14.0, size=n)
preds = [max(0.0, round(g + s, 2))
          for g, s in zip(gt, shifts)]
```

The trained model is loaded and the result discarded. Predictions are ground truth plus noise at a fixed seed, and the noise is tuned—the docstring says so—to land the downstream error metrics inside the range the grader accepts. Success on those scenarios therefore records whether the agent called the tools in a workable order and reported what they returned. It records nothing about prognostic accuracy, which belongs to the generator.

Third, we did not repair the benchmark. Modifying a benchmark's tools forfeits comparability with everything published on it. We instead use the split as a design: 45 scenarios whose tools compute deterministically from what the agent supplies, and 30 whose outcome is fabricated.

IV. A GATE THAT FAILS CLOSED

We release a provenance gate that runs before any modelling and exits non-zero when no measured signal family survives. Four checks correspond to the defects above: constant columns, distributions matching a fallback generator, harness verdict strings in fields features can reach, and columns flagged unmeasured. The synthetic-logprob test is not tuned to one

generator: it keys on three properties real logprobs have—strong negative skew, poor Gaussian fit, and non-trivial mass above -0.01 . Measured on the original corpus these read -0.20 and 4.3% ; on repaired instrumentation, -5.15 and 86.8% .

Two further checks came from failures of the repaired pipeline, and are the ones we would most want others to adopt.

Degenerate trajectories. A run recording no extracted action is a failed collection, not a hard instance. A provider that refuses by returning an empty generation is indistinguishable from a weak model: the agent parses nothing, exhausts its step budget on format errors, and writes an ordinary-looking failure. This destroyed 200 of 225 trajectories in one of our runs, at a tenth of the normal wall time.

Availability leakage. Provenance asks whether each recorded value is a measurement. It cannot see the case where every value is genuine but the rows carrying the column are also the rows that succeed—then a model separates the classes by detecting its own availability, and the imputation filling the gap becomes the signal. In the run above, token features survived on 25 of 225 trajectories and all 11 successes lay inside them; every family scored above 0.9 AUROC while measuring only which runs had been collected. The gate now compares outcome rates on rows that have a column against those that do not.

A self-test builds one honestly captured corpus and one reproducing the defects of Sec. III, and asserts the gate accepts the first and rejects the second.

V. WHAT CHANGES WHEN THE FEATURES ARE REAL

To ask the original question at all we repaired the instrumentation and re-collected. This required leaving the managed inference gateway entirely: it refuses the logprobs parameter on every open-weights model it hosts—verified on five, on both its endpoints, as a warning rather than an error, so the call succeeds and the field is silently absent. Only its commercial model returns them. That is an awkward constraint for a field whose deployments are the least able to use a commercial API, and it means token-level uncertainty is unobservable on gateway-hosted open weights.

We re-ran all 75 scenarios on three open-weights backbones served over an OpenAI-compatible API: 225 trajectories, Pass@1 0.556, one degenerate run, gate-certified. Table II reports the computed-outcome stratum.

Three observations. **Token uncertainty separates on two backbones of three**, at AUROC 0.776 and 0.746, surviving correction over all comparisons made. Out-of-fold permutation importance—preferred to impurity-based measures, which are biased toward high-cardinality continuous predictors [9]—attributes essentially all of it to the entropy of the first two steps; variance over the whole trajectory contributes nothing. Operationally that favours abstaining after two steps rather than executing a full trajectory and discarding it.

The model’s own confidence does not separate (0.494 on the backbone where entropy scores 0.776). We elicited it post hoc, from the model that produced each trajectory,

TABLE II
OUT-OF-FOLD AUROC ON THE 45 SCENARIOS WHOSE OUTCOME THE TOOLS ACTUALLY COMPUTE, PER BACKBONE ($n = 45$ EACH). * SIGNIFICANT AT THE BONFERRONI THRESHOLD OVER ALL 15 COMPARISONS ($\alpha = 0.0033$).

Signal	qwen2.5-7b	llama-3.3-70b	gpt-oss-120b
Early-step entropy	0.776*	0.539	0.746*
Structural	0.648	0.547	0.441
Observation	0.664	0.589	0.312
Self-reported conf.	0.494	0.577	0.421
<i>Pooled, 3 backbones</i> 0.517–0.567, none significant			

because asking in band made models over-generate past the action input, corrupting the tool argument and driving Pass@1 to zero. The elicited values are well dispersed (9 distinct values, $\sigma = 0.36$), so this is a discriminating baseline that discriminates poorly—not a degenerate one.

Pooling destroys the effect. Features identify the generating backbone in 92% of rows, so mixing three populations makes the rankings incommensurable: a logprob that is low for one model is not low in the same way for another. Trajectory calibration appears to be a per-model exercise.

VI. CHECKLIST, LIMITATIONS, AVAILABILITY

Beyond the six detectors, four checks cost little and would have caught most of what we found.

Declare provenance per feature. State for each column whether it is provider-returned, parsed, derived, or imputed, and report the imputation rate. A column imputed on every row is not evidence, whatever its distribution looks like.

Reconcile runs attempted against runs recorded. Count them separately and require them to match. No inspection of values can find a missing row, because the surviving rows are individually valid; a loop that logs only what it wrote reports complete success while discarding cases. An encoding error on agent output cost us 44% of one run, selectively, and the survivors looked healthier than the full set.

Verify the provider answered, not that the call returned. Refusal need not raise. Treat runs that finish implausibly fast as suspect, and probe each backbone for a non-empty generation carrying the fields the study depends on before collecting anything.

Group folds by scenario, and report both trivial policies. Multi-backbone sweeps over a shared scenario set leak across models otherwise. And a selective method must be compared against always-abstain as well as always-execute: in the original study the “optimised” policy was twice as expensive as abstaining on everything, which no reported number revealed because that baseline was never computed.

Limitations. The empirical section illustrates rather than establishes. Each cell holds $n = 45$, with intervals near ± 0.16 ; two of three backbones show the effect and we cannot explain the third; all three are quantised variants from one provider; and the significant cells were found after exploring fifteen comparisons, so they warrant confirmation on fresh runs rather

than treatment as settled. The audit findings in Sec. III carry no such caveat—they are demonstrations, not estimates.

Code, gate, telemetry, and the harness that reproduces every number: <https://github.com/altieriumb7/trajectory-calibration-digital-twins>.

REFERENCES

- [1] Anthropic, “Model Context Protocol specification,” 2024. [Online]. Available: <https://modelcontextprotocol.io>
- [2] T. Feng *et al.*, “PHMForge: Evaluating LLM agents on industrial prognostics through MCP-native, algorithm-grounded tools,” *arXiv:2604.01532*, 2026.
- [3] D. Patel *et al.*, “AssetOpsBench: Benchmarking AI agents for task automation in industrial asset operations and maintenance,” *arXiv:2506.03828*, 2025.
- [4] S. Yao *et al.*, “ReAct: Synergizing reasoning and acting in language models,” in *Proc. ICLR*, 2023.
- [5] S. Kadavath *et al.*, “Language models (mostly) know what they know,” *arXiv:2207.05221*, 2022.
- [6] M. Xiong *et al.*, “Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs,” in *Proc. ICLR*, 2024.
- [7] A. Vinod, Y. Ding, and E. Stengel-Eskin, “CalVerT: Augmenting agents with calibrated verifier telemetry improves action and learning in knowledge-intensive tasks,” *arXiv:2606.21777*, 2026.
- [8] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proc. ICML*, 2017.
- [9] C. Strobl, A.-L. Boulesteix, A. Zeileis, and T. Hothorn, “Bias in random forest variable importance measures,” *BMC Bioinformatics*, vol. 8, no. 25, 2007.
- [10] S. Kapoor and A. Narayanan, “Leakage and the reproducibility crisis in machine-learning-based science,” *Patterns*, vol. 4, no. 9, 2023.